

Attri-VAE for Interpretable Prediction of Myocardial Infarction Complications

Mithun Manivannan

2025-04-23

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, TensorDataset
from torch.optim.lr_scheduler import ReduceLROnPlateau

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, average_precision_score, confusion_matrix
from sklearn.manifold import TSNE
from ucimlrepo import fetch_ucirepo

import os
import json
from pathlib import Path
from datetime import datetime
import logging
from tqdm.notebook import tqdm
import random
from collections import defaultdict
import copy

# Set up logging
logging.basicConfig(level=logging.INFO, format='%(asctime)s - %(levelname)s - %(message)s')
```

```

logger = logging.getLogger()

# Set random seed for reproducibility
def set_seed(seed=42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed(seed)
        torch.cuda.manual_seed_all(seed)
        torch.backends.cudnn.deterministic = True
        torch.backends.cudnn.benchmark = False

set_seed(42)

# Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")

```

Using device: cpu

```

class MIDataset(Dataset):
    """
    Custom dataset for MI complications data with temporal structure
    """
    def __init__(
        self,
        X_timepoints,
        y,
        attributes,
        transform=None,
        target_transform=None
    ):
        """
        Initialize dataset

        Args:
            X_timepoints: Dictionary of feature dataframes for each timepoint
            y: Target array
            attributes: Clinical attributes for regularization
            transform: Optional transform to be applied to features

```

```

        target_transform: Optional transform to be applied to targets
    """
    self.X_timepoints = X_timepoints
    self.y = y
    self.attributes = attributes
    self.transform = transform
    self.target_transform = target_transform

    # Convert to tensors
    self.X_tensors = {
        tp: torch.FloatTensor(X.values)
        for tp, X in X_timepoints.items()
    }
    self.y_tensor = torch.FloatTensor(y.values)
    self.attributes_tensor = torch.FloatTensor(attributes.values)

def __len__(self):
    return len(self.y)

def __getitem__(self, idx):
    # Get features for each timepoint
    x_list = [self.X_tensors[tp][idx] for tp in sorted(self.X_tensors.keys())]

    # Get target and attributes
    y = self.y_tensor[idx]
    attributes = self.attributes_tensor[idx]

    # Apply transforms if specified
    if self.transform:
        x_list = [self.transform(x) for x in x_list]
    if self.target_transform:
        y = self.target_transform(y)

    return x_list, y, attributes

def load_mi_data():
    """
    Load and preprocess MI complications dataset

    Returns:
        X_timepoints: Dictionary of feature dataframes for each timepoint
        y: Target array
    """

```

```

        attributes: Clinical attributes for regularization
"""
print("Loading MI Complications dataset from UCI repository...")
# Fetch the UCI MI complications dataset
mi = fetch_ucirepo(id=579)
X_full = mi.data.features
y_full = mi.data.targets

# Print dataset information
print(f"Dataset shape: {X_full.shape}")
print(f"Number of samples: {len(X_full)}")
print(f"Number of complications: {y_full.shape[1]}")

# Checking if dataset is loaded properly
if X_full is None or X_full.empty:
    raise ValueError("Failed to load features from UCI MI complications dataset")

# Create any-complication binary target
y = pd.Series((y_full.sum(axis=1) > 0).astype(int))
print(f"Positive samples (with complications): {y.sum()} ({y.mean()*100:.2f}%)")

# Define timepoints and their corresponding features
timepoints = {
    'admission': [col for col in X_full.columns if '_ADM' in col or col in ['AGE', 'SEX',
    '24h': [col for col in X_full.columns if '_24_' in col or '_24H' in col],
    '48h': [col for col in X_full.columns if '_48_' in col or '_48H' in col],
    '72h': [col for col in X_full.columns if '_72_' in col or '_72H' in col]
}

# Print timepoint information
for tp, cols in timepoints.items():
    print(f"Timepoint {tp}: {len(cols)} features")

# Create feature sets for each timepoint
X_timepoints = {}
for tp, cols in timepoints.items():
    if cols: # Only process if we have columns for this timepoint
        X_tp = X_full[cols].copy()
        X_timepoints[tp] = X_tp

# Select clinical attributes for regularization
# These should be features that are clinically meaningful

```

```

attribute_cols = [
    'AGE',
    'S_AD_ORIT',
    'D_AD_ORIT',
    'K_BLOOD',
    'L_BLOOD',
    'TIME_B_S'
]

# Filter to only include columns that exist in the dataset
attribute_cols = [col for col in attribute_cols if col in X_full.columns]
attributes = X_full[attribute_cols].copy()
print(f"Selected {len(attribute_cols)} clinical attributes: {attribute_cols}")

# Handle missing values
for tp, X in X_timepoints.items():
    # Check if any columns are fully NaN
    empty_cols = X.columns[X.isna().all()].tolist()
    if empty_cols:
        print(f"Warning: Dropping columns with all NaN values in {tp}: {empty_cols}")
        X.drop(columns=empty_cols, inplace=True)

    # Calculate percentage of missing values
    missing_pct = X.isna().mean().mean() * 100
    print(f"Timepoint {tp}: {missing_pct:.2f}% missing values")

    # Now fill remaining NaNs with column means
    X.fillna(X.mean(), inplace=True)

attributes.fillna(attributes.mean(), inplace=True)

# Scale features
scalers = {}
for tp, X in X_timepoints.items():
    if not X.empty: # Only scale if we have data
        scaler = StandardScaler()
        X_timepoints[tp] = pd.DataFrame(
            scaler.fit_transform(X),
            columns=X.columns,
            index=X.index
        )
        scalers[tp] = scaler

```

```

# Scale attributes
if not attributes.empty:
    attr_scaler = StandardScaler()
    attributes = pd.DataFrame(
        attr_scaler.fit_transform(attributes),
        columns=attributes.columns,
        index=attributes.index
    )

print("Data preprocessing completed successfully.")
return X_timepoints, y, attributes

def create_data_loaders(
    X_timepoints,
    y,
    attributes,
    batch_size=32,
    val_size=0.2,
    test_size=0.1,
    random_state=42
):
    """
    Create train, validation, and test data loaders

    Args:
        X_timepoints: Dictionary of feature dataframes for each timepoint
        y: Target array
        attributes: Clinical attributes for regularization
        batch_size: Batch size for data loaders
        val_size: Proportion of data to use for validation
        test_size: Proportion of data to use for testing
        random_state: Random seed for reproducibility

    Returns:
        train_loader: DataLoader for training data
        val_loader: DataLoader for validation data
        test_loader: DataLoader for test data
    """
    print("Creating train/val/test splits...")
    # First split into train+val and test
    train_val_idx, test_idx = train_test_split(
        np.arange(len(y)),

```

```

        test_size=test_size,
        random_state=random_state,
        stratify=y
    )

    # Then split train+val into train and val
    train_idx, val_idx = train_test_split(
        train_val_idx,
        test_size=val_size/(1-test_size),
        random_state=random_state,
        stratify=y.iloc[train_val_idx]
    )

    print(f"Train set: {len(train_idx)} samples ({y.iloc[train_idx].mean()*100:.2f}% positive)
    print(f"Validation set: {len(val_idx)} samples ({y.iloc[val_idx].mean()*100:.2f}% positive)
    print(f"Test set: {len(test_idx)} samples ({y.iloc[test_idx].mean()*100:.2f}% positive)".

    # Create datasets
    train_dataset = MIDataset(
        {tp: X.iloc[train_idx] for tp, X in X_timepoints.items()},
        y.iloc[train_idx],
        attributes.iloc[train_idx]
    )

    val_dataset = MIDataset(
        {tp: X.iloc[val_idx] for tp, X in X_timepoints.items()},
        y.iloc[val_idx],
        attributes.iloc[val_idx]
    )

    test_dataset = MIDataset(
        {tp: X.iloc[test_idx] for tp, X in X_timepoints.items()},
        y.iloc[test_idx],
        attributes.iloc[test_idx]
    )

    # Create data loaders with num_workers=0 to avoid multiprocessing issues in Jupyter
    train_loader = DataLoader(
        train_dataset,
        batch_size=batch_size,
        shuffle=True,
        num_workers=0 # Changed from 4 to 0

```

```

    )

    val_loader = DataLoader(
        val_dataset,
        batch_size=batch_size,
        shuffle=False,
        num_workers=0 # Changed from 4 to 0
    )

    test_loader = DataLoader(
        test_dataset,
        batch_size=batch_size,
        shuffle=False,
        num_workers=0 # Changed from 4 to 0
    )

    print("Data loaders created successfully.")
    return train_loader, val_loader, test_loader

# Load and prepare data
X_timepoints, y, attributes = load_mi_data()
train_loader, val_loader, test_loader = create_data_loaders(X_timepoints, y, attributes)

# Extract feature dimensions for model configuration
input_dims = [X.shape[1] for X in X_timepoints.values()]
attribute_dims = {attr: i for i, attr in enumerate(attributes.columns)}
print(f"Input dimensions: {input_dims}")
print(f"Attribute dimensions: {attribute_dims}")

```

Loading MI Complications dataset from UCI repository...

Dataset shape: (1700, 111)

Number of samples: 1700

Number of complications: 12

Positive samples (with complications): 1037 (61.00%)

Timepoint admission: 3 features

Timepoint 24h: 0 features

Timepoint 48h: 0 features

Timepoint 72h: 0 features

Selected 6 clinical attributes: ['AGE', 'S_AD_ORIT', 'D_AD_ORIT', 'K_BLOOD', 'L_BLOOD', 'TIM

Timepoint admission: 2.63% missing values

Data preprocessing completed successfully.

Creating train/val/test splits...

Train set: 1190 samples (61.01% positive)
 Validation set: 340 samples (60.88% positive)
 Test set: 170 samples (61.18% positive)
 Data loaders created successfully.
 Input dimensions: [3]
 Attribute dimensions: {'AGE': 0, 'S_AD_ORIT': 1, 'D_AD_ORIT': 2, 'K_BLOOD': 3, 'L_BLOOD': 4,

```

class MedicalSafetyLayer(nn.Module):
    """
    Layer to enforce physiological ranges for outputs
    """
    def __init__(self, feature_ranges=None):
        """
        Initialize layer with physiological ranges

        Args:
            feature_ranges: Dictionary mapping feature names to (min, max) range tuples
        """
        super(MedicalSafetyLayer, self).__init__()
        self.feature_ranges = feature_ranges or {
            # Systolic BP (mmHg)
            'S_AD': (80, 180),
            # Diastolic BP (mmHg)
            'D_AD': (40, 120),
            # Heart rate (bpm)
            'RATE_AD': (40, 180),
            # Body temperature (°C)
            'TEMP_AD': (35, 41),
            # Potassium (mmol/L)
            'K_BLOOD': (3.0, 6.0),
            # Sodium (mmol/L)
            'Na_BLOOD': (130, 150),
            # White blood cells (109/L)
            'L_BLOOD': (4.0, 25.0)
        }

    def forward(self, x, feature_names=None):
        """
        Apply safety constraints to outputs

        Args:
            x: Input tensor
  
```

```

        feature_names: List of feature names corresponding to dimensions

Returns:
    x_safe: Output tensor with values within safe ranges
    safety_loss: Loss penalizing values outside physiological ranges
"""
# If feature names not provided, return unchanged
if feature_names is None:
    return x, None

# Initialize loss
safety_loss = 0.0

# Apply safety constraints
for i, name in enumerate(feature_names):
    for key, (min_val, max_val) in self.feature_ranges.items():
        if key in name:
            # Calculate penalty for out-of-range values
            below_min = F.relu(min_val - x[:, i])
            above_max = F.relu(x[:, i] - max_val)
            penalty = below_min + above_max
            safety_loss += torch.mean(penalty)

    return x, safety_loss

class GroupLatentToTimepoint(nn.Module):
    """
    Group-sparse latent-to-timepoint transformation as in oi-VAE
    """
    def __init__(self, latent_dim, hidden_dim, timepoints):
        super(GroupLatentToTimepoint, self).__init__()
        self.latent_dim = latent_dim
        self.hidden_dim = hidden_dim
        self.timepoints = timepoints

        # Group-specific latent-to-timepoint matrices (one per timepoint)
        # These will be subject to group sparsity regularization
        self.W = nn.ParameterDict({
            tp: nn.Parameter(torch.randn(hidden_dim, latent_dim) * 0.01)
            for tp in timepoints
        })

```

```

def forward(self, z):
    """
    Transform latent vector with group-specific matrices

    Args:
        z: Latent vector [batch_size, latent_dim]

    Returns:
        outputs: Dictionary of timepoint-specific transformed vectors
    """
    outputs = {}
    for tp in self.timepoints:
        # Apply group-specific transformation
        outputs[tp] = F.linear(z, self.W[tp])
    return outputs

def get_group_norms(self):
    """
    Compute the L2 norm of each column in each group matrix
    Used for proximal gradient updates and visualization

    Returns:
        Dictionary of norms per timepoint and latent dimension
    """
    norms = {}
    for tp in self.timepoints:
        # Compute column-wise L2 norms: [latent_dim]
        norms[tp] = torch.norm(self.W[tp], dim=0)
    return norms

def apply_proximal_update(self, lr, lambda_reg):
    """
    Apply proximal gradient update for group sparsity

    Args:
        lr: Learning rate
        lambda_reg: Regularization strength
    """
    with torch.no_grad():
        for tp in self.timepoints:
            # Compute column-wise L2 norms
            norms = torch.norm(self.W[tp], dim=0, keepdim=True)

```

```

        # Apply proximal operator: shrink columns toward zero
        # max(0, 1 - lambda*lr/norm) * w
        scale = torch.clamp(1 - lambda_reg * lr / (norms + 1e-8), min=0.0)
        self.W[tp].mul_(scale)

class GroupInterpretableVAE(nn.Module):
    """
    Group-structured interpretable VAE for MI complications prediction with attribute regularization
    """
    def __init__(
        self,
        input_dims,
        timepoint_names,
        latent_dim=32,
        embed_dim=16,
        hidden_dim=64,
        attribute_dims=None,
        medical_safety=True
    ):
        """
        Initialize model

        Args:
            input_dims: List of input dimensions for each timepoint
            timepoint_names: List of timepoint names (e.g., 'admission', '24h')
            latent_dim: Dimension of the latent space
            embed_dim: Dimension of time embeddings
            hidden_dim: Dimension of hidden layers
            attribute_dims: Dictionary mapping attribute name to latent dimension index
            medical_safety: Whether to enforce medical safety constraints
        """
        super(GroupInterpretableVAE, self).__init__()

        self.input_dims = input_dims if isinstance(input_dims, list) else list(input_dims)
        self.timepoint_names = timepoint_names
        self.latent_dim = latent_dim
        self.embed_dim = embed_dim
        self.hidden_dim = hidden_dim
        self.attribute_dims = attribute_dims or {}
        self.medical_safety = medical_safety

        # Time embeddings (one for each timepoint)

```

```

self.time_embeddings = nn.Embedding(len(input_dims), embed_dim)

# Encoders for each timepoint
self.encoders = nn.ModuleList([
    nn.Sequential(
        nn.Linear(dim + embed_dim, hidden_dim),
        nn.BatchNorm1d(hidden_dim),
        nn.ReLU(),
        nn.Linear(hidden_dim, hidden_dim // 2),
        nn.BatchNorm1d(hidden_dim // 2),
        nn.ReLU()
    ) for dim in input_dims
])

# Latent projectors (mu and logvar)
encoder_output_dim = (hidden_dim // 2) * len(input_dims)
self.mu_projector = nn.Linear(encoder_output_dim, latent_dim)
self.logvar_projector = nn.Linear(encoder_output_dim, latent_dim)

# Group-specific latent-to-timepoint transformation with sparsity
self.group_transform = GroupLatentToTimepoint(latent_dim, hidden_dim, timepoint_names)

# Integration module
self.integration = nn.Sequential(
    nn.Linear(latent_dim, latent_dim),
    nn.BatchNorm1d(latent_dim),
    nn.ReLU()
)

# Attribute predictor for regularization
self.attribute_predictor = nn.Linear(latent_dim, len(attribute_dims))

# Classifier for complications prediction
self.classifier = nn.Sequential(
    nn.Linear(latent_dim, hidden_dim // 2),
    nn.BatchNorm1d(hidden_dim // 2),
    nn.ReLU(),
    nn.Dropout(0.2),
    nn.Linear(hidden_dim // 2, hidden_dim // 4),
    nn.BatchNorm1d(hidden_dim // 4),
    nn.ReLU(),
    nn.Dropout(0.2),

```

```

        nn.Linear(hidden_dim // 4, 1)
    )

    # Decoders for each timepoint
    self.decoders = nn.ModuleDict({
        tp: nn.Sequential(
            nn.Linear(hidden_dim + embed_dim, hidden_dim),
            nn.BatchNorm1d(hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, dim)
        ) for tp, dim in zip(timepoint_names, input_dims)
    })

    # Medical safety layer
    if medical_safety:
        self.safety_layer = MedicalSafetyLayer()

    # Loss functions
    self.mse_loss = nn.MSELoss(reduction='mean')
    self.bce_loss = nn.BCEWithLogitsLoss(reduction='mean')

def encode(self, x_list, attributes=None):
    """
    Encode a list of inputs from different timepoints

    Args:
        x_list: List of input tensors
        attributes: Clinical attributes tensor

    Returns:
        mu: Mean vector
        logvar: Log variance vector
    """
    batch_size = x_list[0].size(0)
    encodings = []

    # Generate time embeddings
    time_indices = torch.arange(len(self.encoders), device=x_list[0].device)
    time_embeds = self.time_embeddings(time_indices)

    # Encode each timepoint with its time embedding
    for i, (x, encoder) in enumerate(zip(x_list, self.encoders)):

```

```

        # Add time embedding to input
        time_embed = time_embeds[i].expand(batch_size, -1)
        x_t = torch.cat([x, time_embed], dim=1)

        # Encode
        h = encoder(x_t)
        encodings.append(h)

    # Concatenate all encodings
    concat_encoding = torch.cat(encodings, dim=1)

    # Project to latent space
    mu = self.mu_projector(concat_encoding)
    logvar = self.logvar_projector(concat_encoding)

    return mu, logvar

def reparameterize(self, mu, logvar):
    """
    Reparameterization trick to sample from distribution

    Args:
        mu: Mean vector
        logvar: Log variance vector

    Returns:
        z: Sampled latent vector
    """
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)
    return mu + eps * std

def decode(self, z):
    """
    Decode latent representation using group-sparse transformations

    Args:
        z: Latent vector

    Returns:
        recon_list: List of reconstructions for each timepoint
        safety_loss: Medical safety loss if enabled

```

```

"""
batch_size = z.size(0)
recon_list = []
safety_loss = None

# Apply group-specific latent-to-timepoint transformations with sparsity
transformed = self.group_transform(z)

# Generate time embeddings
time_indices = torch.arange(len(self.timepoint_names), device=z.device)
time_embeds = self.time_embeddings(time_indices)

# Decode each timepoint
for i, tp in enumerate(self.timepoint_names):
    # Add time embedding
    time_embed = time_embeds[i].expand(batch_size, -1)
    h = torch.cat([transformed[tp], time_embed], dim=1)

    # Decode
    x_recon = self.decoders[tp](h)

    # Apply medical safety constraints if enabled
    if self.medical_safety:
        x_recon, time_safety_loss = self.safety_layer(x_recon)
        if time_safety_loss is not None:
            if safety_loss is None:
                safety_loss = time_safety_loss
            else:
                safety_loss += time_safety_loss

    recon_list.append(x_recon)

return recon_list, safety_loss

def predict_attributes(self, z):
    """
    Predict attributes from latent vector for regularization

    Args:
        z: Latent vector

    Returns:

```



```

        attr_pred: Predicted attributes
    """
    return self.attribute_predictor(z)

def forward(self, x_list, attributes=None):
    """
    Forward pass

    Args:
        x_list: List of input tensors
        attributes: Clinical attributes tensor

    Returns:
        recon_list: List of reconstructions
        mu: Mean vector
        logvar: Log variance vector
        attr_pred: Predicted attributes
        y_pred: Predicted complication probability
        safety_loss: Medical safety loss if enabled
    """
    # Encode
    mu, logvar = self.encode(x_list, attributes)

    # Sample latent vector
    z = self.reparameterize(mu, logvar)

    # Integrate latent representation
    z = self.integration(z)

    # Predict attributes
    attr_pred = self.predict_attributes(z)

    # Predict complications
    y_pred = self.classifier(z)

    # Decode
    recon_list, safety_loss = self.decode(z)

    return recon_list, mu, logvar, attr_pred, y_pred, safety_loss

def compute_losses(self, outputs, targets, attributes):
    """

```

Compute all loss components

Args:

outputs: Model outputs (recon_list, mu, logvar, attr_pred, y_pred, safety_loss)
targets: Target values (x_list, y, _)
attributes: Clinical attributes

Returns:

total_loss: Total weighted loss
loss_dict: Dictionary of individual loss components

"""

Unpack outputs

reconstructions, mu, logvar, pred_attributes, y_pred, safety_loss = outputs
x_list, y, _ = targets

Ensure y and y_pred have the same shape

y = y.view(-1, 1) if y_pred.shape[-1] == 1 else y

1. Reconstruction loss (MSE)

recon_loss = 0

for i, (x, recon) in enumerate(zip(x_list, reconstructions)):

recon_loss += self.mse_loss(recon, x)

recon_loss /= len(x_list)

2. KL divergence loss

kl_loss = -0.5 * torch.mean(1 + logvar - mu.pow(2) - logvar.exp())

3. Attribute regularization loss

attr_loss = self.mse_loss(pred_attributes, attributes)

4. Classification loss for complications

cls_loss = self.bce_loss(y_pred, y)

5. Group sparsity loss (computed during optimization)

Total loss with weighting

total_loss = recon_loss + self.beta * kl_loss + self.gamma * attr_loss + cls_loss

Add medical safety loss if available

if safety_loss is not None:

total_loss += self.delta * safety_loss

```

# Return individual losses for logging
loss_dict = {
    'loss': total_loss.item(),
    'recon_loss': recon_loss.item(),
    'kl_loss': kl_loss.item(),
    'attr_loss': attr_loss.item(),
    'cls_loss': cls_loss.item(),
}

if safety_loss is not None:
    loss_dict['safety_loss'] = safety_loss.item()

return total_loss, loss_dict

def apply_proximal_update(self, lr, lambda_reg):
    """
    Apply proximal gradient update for group sparsity

    Args:
        lr: Learning rate
        lambda_reg: Regularization strength
    """
    self.group_transform.apply_proximal_update(lr, lambda_reg)

def get_group_sparsity_visualization(self):
    """
    Get matrix of group norms for visualization

    Returns:
        Group norm matrix for heatmap visualization
    """
    norms = self.group_transform.get_group_norms()

    # Convert to numpy for visualization
    norm_matrix = np.zeros((len(self.timepoint_names), self.latent_dim))
    for i, tp in enumerate(self.timepoint_names):
        norm_matrix[i] = norms[tp].cpu().detach().numpy()

    return norm_matrix

```

```

class TensorboardCallback:
    """Callback for logging metrics to TensorBoard during training"""

```

```

def __init__(self, writer):
    self.writer = writer
    self.step = 0

def on_epoch_end(self, epoch, train_metrics, val_metrics):
    """Log metrics at the end of each epoch"""
    # Log training metrics
    for k, v in train_metrics.items():
        if isinstance(v, (int, float)):
            self.writer.add_scalar(f'train/{k}', v, epoch)

    # Log validation metrics
    for k, v in val_metrics.items():
        if isinstance(v, (int, float)):
            self.writer.add_scalar(f'val/{k}', v, epoch)

    # Log learning rate
    for i, param_group in enumerate(self.optimizer.param_groups):
        self.writer.add_scalar(f'lr/group_{i}', param_group['lr'], epoch)

def set_optimizer(self, optimizer):
    """Set the optimizer for learning rate tracking"""
    self.optimizer = optimizer

class GroupInterpretableVAETrainer:
    """
    Trainer for Group Interpretable VAE with collapsed variational inference
    """
    def __init__(
        self,
        model,
        train_loader,
        val_loader,
        device,
        lr=1e-3,
        weight_decay=1e-5,
        beta=1.0,          # KL weight
        gamma=0.1,         # Attribute regularization weight
        delta=0.1,         # Medical safety weight
        lambda_reg=1.0,    # Group sparsity weight
        callbacks=None
    ):

```

```

"""
Initialize trainer

Args:
    model: GroupInterpretableVAE model
    train_loader: DataLoader for training data
    val_loader: DataLoader for validation data
    device: Device to use (cuda or cpu)
    lr: Learning rate
    weight_decay: Weight decay
    beta: KL divergence weight
    gamma: Attribute regularization weight
    delta: Medical safety weight
    lambda_reg: Group sparsity regularization strength
    callbacks: List of callbacks for training events
"""

self.model = model.to(device)
self.train_loader = train_loader
self.val_loader = val_loader
self.device = device
self.lr = lr
self.lambda_reg = lambda_reg

# Set loss weights
self.model.beta = beta
self.model.gamma = gamma
self.model.delta = delta

# Initialize optimizer
self.optimizer = optim.Adam(
    model.parameters(),
    lr=lr,
    weight_decay=weight_decay
)

# Initialize learning rate scheduler
self.scheduler = ReduceLROnPlateau(
    self.optimizer,
    mode='min',
    factor=0.5,
    patience=5
)

```

```

# Callbacks
self.callbacks = callbacks or []

# Set optimizer in callbacks
for callback in self.callbacks:
    if hasattr(callback, 'set_optimizer'):
        callback.set_optimizer(self.optimizer)

def train_epoch(self, epoch):
    """
    Train model for one epoch with collapsed variational inference

    Args:
        epoch: Current epoch number

    Returns:
        avg_loss: Average loss for the epoch
        metrics: Dictionary of metrics
    """
    self.model.train()
    running_loss = 0.0
    loss_dict_sum = defaultdict(float)
    predictions = []
    targets = []

    # Create progress bar
    pbar = tqdm(self.train_loader, desc=f'Epoch {epoch+1} [Train]')

    for batch_idx, (x_list, y, attributes) in enumerate(pbar):
        # Move data to device
        x_list = [x.to(self.device) for x in x_list]
        y = y.to(self.device)
        attributes = attributes.to(self.device)

        # Zero gradients
        self.optimizer.zero_grad()

        # Forward pass
        outputs = self.model(x_list, attributes)

        # Compute loss
        loss, batch_loss_dict = self.model.compute_losses(outputs, (x_list, y, attributes)

```

```

        # Backward pass
        loss.backward()

        # Update weights with standard gradient descent
        self.optimizer.step()

        # Apply proximal update for group sparsity
        self.model.apply_proximal_update(self.lr, self.lambda_reg)

        # Update running loss
        running_loss += loss.item()
        for k, v in batch_loss_dict.items():
            loss_dict_sum[k] += v

        # Update progress bar
        pbar.set_postfix({'loss': loss.item()})

        # Store predictions for metrics
        y_pred = outputs[4]
        predictions.append(y_pred.detach().cpu().numpy())
        targets.append(y.detach().cpu().numpy())

    # Concatenate predictions and targets
    predictions = np.concatenate(predictions)
    targets = np.concatenate(targets)

    # Reshape for metrics calculation
    predictions = predictions.reshape(-1)
    targets = targets.reshape(-1)

    # Calculate metrics
    metrics = {}
    for k, v in loss_dict_sum.items():
        metrics[k] = v / len(self.train_loader)

    # Calculate classification metrics
    try:
        metrics['auroc'] = roc_auc_score(targets, predictions)
        metrics['auprc'] = average_precision_score(targets, predictions)
    except Exception as e:
        print(f"Error calculating metrics: {e}")
        metrics['auroc'] = 0.0

```

```

        metrics['auprc'] = 0.0

    return metrics['loss'], metrics

def validate(self, dataloader=None):
    """
    Validate model on validation set

    Args:
        dataloader: DataLoader to use for validation (default: self.val_loader)

    Returns:
        avg_loss: Average loss for validation
        metrics: Dictionary of metrics
    """
    if dataloader is None:
        dataloader = self.val_loader

    self.model.eval()
    running_loss = 0.0
    loss_dict_sum = defaultdict(float)
    predictions = []
    targets = []

    with torch.no_grad():
        # Create progress bar
        pbar = tqdm(dataloader, desc='Validation')

        for batch_idx, (x_list, y, attributes) in enumerate(pbar):
            # Move data to device
            x_list = [x.to(self.device) for x in x_list]
            y = y.to(self.device)
            attributes = attributes.to(self.device)

            # Forward pass
            outputs = self.model(x_list, attributes)

            # Compute loss
            loss, batch_loss_dict = self.model.compute_losses(outputs, (x_list, y, attributes))

            # Update running loss
            running_loss += loss.item()

```



```

        for k, v in batch_loss_dict.items():
            loss_dict_sum[k] += v

        # Update progress bar
        pbar.set_postfix({'loss': loss.item()})

        # Store predictions for metrics
        y_pred = outputs[4]
        predictions.append(y_pred.detach().cpu().numpy())
        targets.append(y.detach().cpu().numpy())

    # Concatenate predictions and targets
    predictions = np.concatenate(predictions)
    targets = np.concatenate(targets)

    # Reshape for metrics calculation
    predictions = predictions.reshape(-1)
    targets = targets.reshape(-1)

    # Calculate metrics
    metrics = {}
    for k, v in loss_dict_sum.items():
        metrics[k] = v / len(dataloader)

    # Calculate classification metrics
    try:
        metrics['auroc'] = roc_auc_score(targets, predictions)
        metrics['auprc'] = average_precision_score(targets, predictions)
    except Exception as e:
        print(f"Error calculating metrics: {e}")
        metrics['auroc'] = 0.0
        metrics['auprc'] = 0.0

    return metrics

def train(self, n_epochs=100, early_stopping_patience=10, checkpoint_dir="./checkpoints")
    """
    Train the model with early stopping

    Args:
        n_epochs: Maximum number of epochs to train
        early_stopping_patience: Number of epochs to wait for improvement before stopping

```

```

        checkpoint_dir: Directory to save model checkpoints

Returns:
    history: Dictionary of training and validation metrics
    best_model: Best model state dict
"""
# Create checkpoint directory if it doesn't exist
os.makedirs(checkpoint_dir, exist_ok=True)

# Initialize variables for training
best_val_loss = float('inf')
best_model_state = None
patience_counter = 0
history = {'train': [], 'val': []}

print(f"Starting training for {n_epochs} epochs with early stopping patience {early_

for epoch in range(n_epochs):
    # Train for one epoch
    train_loss, train_metrics = self.train_epoch(epoch)

    # Validate
    val_metrics = self.validate()
    val_loss = val_metrics['loss']

    # Update learning rate scheduler
    self.scheduler.step(val_loss)

    # Save metrics in history
    history['train'].append(train_metrics)
    history['val'].append(val_metrics)

    # Print epoch summary
    print(f"Epoch {epoch+1}/{n_epochs} - "
          f"Train Loss: {train_loss:.4f}, "
          f"Val Loss: {val_loss:.4f}, "
          f"Train AUROC: {train_metrics.get('auroc', 0):.4f}, "
          f"Val AUROC: {val_metrics.get('auroc', 0):.4f}")

    # Call callbacks
    for callback in self.callbacks:
        if hasattr(callback, 'on_epoch_end'):

```

```

        callback.on_epoch_end(epoch, train_metrics, val_metrics)

    # Check for improvement
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_model_state = copy.deepcopy(self.model.state_dict())
        patience_counter = 0

    # Save the best model
    checkpoint_path = os.path.join(checkpoint_dir, f"best_model.pt")
    torch.save({
        'epoch': epoch,
        'model_state_dict': self.model.state_dict(),
        'optimizer_state_dict': self.optimizer.state_dict(),
        'train_metrics': train_metrics,
        'val_metrics': val_metrics,
    }, checkpoint_path)
    print(f"Saved best model at epoch {epoch+1} with val_loss: {val_loss:.4f}")
else:
    patience_counter += 1
    if patience_counter >= early_stopping_patience:
        print(f"Early stopping triggered after {epoch+1} epochs")
        break

# Load the best model
if best_model_state is not None:
    self.model.load_state_dict(best_model_state)

return history, self.model

```

```

def plot_training_history(history):
    """
    Plot training and validation metrics

    Args:
        history: Dictionary of training and validation metrics from trainer
    """
    # Convert history to DataFrame for easier plotting
    epochs = range(1, len(history['train']) + 1)
    metrics = ['loss', 'recon_loss', 'kl_loss', 'attr_loss', 'cls_loss', 'auROC', 'auprc']

    plt.figure(figsize=(15, 12))

```

```

for i, metric in enumerate(metrics):
    if metric not in history['train'][0]:
        continue

    plt.subplot(3, 3, i+1)

    # Get train and val values for the metric
    train_values = [epoch_metrics.get(metric, 0) for epoch_metrics in history['train']]
    val_values = [epoch_metrics.get(metric, 0) for epoch_metrics in history['val']]

    plt.plot(epochs, train_values, 'b-', label=f'Train {metric}')
    plt.plot(epochs, val_values, 'r-', label=f'Val {metric}')

    plt.title(f'{metric} vs. Epochs')
    plt.xlabel('Epochs')
    plt.ylabel(metric)
    plt.legend()

    # Highlight best epoch (lowest val loss or highest val auroc)
    if metric == 'loss':
        best_epoch = np.argmin(val_values) + 1
        plt.axvline(x=best_epoch, color='g', linestyle='--')
        plt.text(best_epoch + 0.1, min(train_values), f'Best epoch: {best_epoch}')
    elif metric == 'auroc':
        best_epoch = np.argmax(val_values) + 1
        plt.axvline(x=best_epoch, color='g', linestyle='--')
        plt.text(best_epoch + 0.1, max(val_values), f'Best epoch: {best_epoch}')

plt.tight_layout()
plt.show()

def visualize_latent_space(model, test_loader, device, plot_type='tsne'):
    """
    Visualize the latent space using t-SNE or PCA

    Args:
        model: Trained model
        test_loader: DataLoader for test data
        device: Device to use
        plot_type: 'tsne' or 'pca'
    """
    model.eval()

```

```

latent_vectors = []
labels = []

with torch.no_grad():
    for x_list, y, attributes in test_loader:
        # Move data to device
        x_list = [x.to(device) for x in x_list]

        # Get latent representation (mu)
        mu, _ = model.encode(x_list)

        # Store latent vectors and labels
        latent_vectors.append(mu.cpu().numpy())
        labels.append(y.cpu().numpy())

# Concatenate batches
latent_vectors = np.concatenate(latent_vectors, axis=0)
labels = np.concatenate(labels, axis=0)

# Apply dimensionality reduction
if plot_type == 'tsne':
    reducer = TSNE(n_components=2, random_state=42)
    reduced_data = reducer.fit_transform(latent_vectors)
    title = 't-SNE Visualization of Latent Space'
else: # pca
    from sklearn.decomposition import PCA
    reducer = PCA(n_components=2, random_state=42)
    reduced_data = reducer.fit_transform(latent_vectors)
    title = 'PCA Visualization of Latent Space'

# Plot
plt.figure(figsize=(10, 8))
scatter = plt.scatter(reduced_data[:, 0], reduced_data[:, 1], c=labels, alpha=0.6,
                     cmap='coolwarm', edgecolors='w', s=100)

plt.colorbar(scatter, label='Complication')
plt.title(title)
plt.xlabel('Dimension 1')
plt.ylabel('Dimension 2')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

```

def visualize_group_sparsity(model):
    """
    Visualize the group sparsity patterns

    Args:
        model: Trained model
    """
    # Get group sparsity matrix
    norm_matrix = model.get_group_sparsity_visualization()

    plt.figure(figsize=(12, 8))
    sns.heatmap(norm_matrix, cmap='viridis',
                xticklabels=[f'z{i+1}' for i in range(model.latent_dim)],
                yticklabels=model.timepoint_names,
                annot=False, cbar=True)

    plt.title('Group Sparsity Pattern')
    plt.xlabel('Latent Dimensions')
    plt.ylabel('Timepoints')
    plt.tight_layout()
    plt.show()

    # Analyze which latent dimensions are active for each timepoint
    for i, tp in enumerate(model.timepoint_names):
        active_dims = np.where(norm_matrix[i] > 0.1)[0] # Threshold can be adjusted
        print(f"Timepoint {tp} is influenced by latent dimensions: {list(active_dims + 1)}")

def visualize_attribute_regularization(model, test_loader, device):
    """
    Visualize how latent dimensions correspond to clinical attributes

    Args:
        model: Trained model
        test_loader: DataLoader for test data
        device: Device to use
    """
    model.eval()
    latent_vectors = []
    all_attributes = []

    with torch.no_grad():
        for x_list, _, attributes in test_loader:

```

```

# Move data to device
x_list = [x.to(device) for x in x_list]

# Get latent representation (mu)
mu, _ = model.encode(x_list)

# Store latent vectors and attributes
latent_vectors.append(mu.cpu().numpy())
all_attributes.append(attributes.cpu().numpy())

# Concatenate batches
latent_vectors = np.concatenate(latent_vectors, axis=0)
all_attributes = np.concatenate(all_attributes, axis=0)

# Calculate correlations between latent dimensions and attributes
corr_matrix = np.zeros((model.latent_dim, all_attributes.shape[1]))

for i in range(model.latent_dim):
    for j in range(all_attributes.shape[1]):
        corr = np.corrcoef(latent_vectors[:, i], all_attributes[:, j])[0, 1]
        corr_matrix[i, j] = corr

# Plot correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, cmap='coolwarm', center=0,
            xticklabels=list(model.attribute_dims.keys()),
            yticklabels=[f'z{i+1}' for i in range(model.latent_dim)],
            annot=True, fmt='.2f')

plt.title('Correlation between Latent Dimensions and Clinical Attributes')
plt.xlabel('Clinical Attributes')

```

```

def plot_training_history(history):
    """
    Plot training and validation metrics

    Args:
        history: Dictionary of training and validation metrics from trainer
    """
    # Convert history to DataFrame for easier plotting
    epochs = range(1, len(history['train']) + 1)
    metrics = ['loss', 'recon_loss', 'kl_loss', 'attr_loss', 'cls_loss', 'auroc', 'auprc']

```

```

plt.figure(figsize=(15, 12))

for i, metric in enumerate(metrics):
    if metric not in history['train'][0]:
        continue

    plt.subplot(3, 3, i+1)

    # Get train and val values for the metric
    train_values = [epoch_metrics.get(metric, 0) for epoch_metrics in history['train']]
    val_values = [epoch_metrics.get(metric, 0) for epoch_metrics in history['val']]

    plt.plot(epochs, train_values, 'b-', label=f'Train {metric}')
    plt.plot(epochs, val_values, 'r-', label=f'Val {metric}')

    plt.title(f'{metric} vs. Epochs')
    plt.xlabel('Epochs')
    plt.ylabel(metric)
    plt.legend()

    # Highlight best epoch (lowest val loss or highest val auROC)
    if metric == 'loss':
        best_epoch = np.argmin(val_values) + 1
        plt.axvline(x=best_epoch, color='g', linestyle='--')
        plt.text(best_epoch + 0.1, min(train_values), f'Best epoch: {best_epoch}')
    elif metric == 'auROC':
        best_epoch = np.argmax(val_values) + 1
        plt.axvline(x=best_epoch, color='g', linestyle='--')
        plt.text(best_epoch + 0.1, max(val_values), f'Best epoch: {best_epoch}')

plt.tight_layout()
plt.show()

def visualize_latent_space(model, test_loader, device, plot_type='tsne'):
    """
    Visualize the latent space using t-SNE or PCA

    Args:
        model: Trained model
        test_loader: DataLoader for test data
        device: Device to use
        plot_type: 'tsne' or 'pca'

```



```

"""
model.eval()
latent_vectors = []
labels = []

with torch.no_grad():
    for x_list, y, attributes in test_loader:
        # Move data to device
        x_list = [x.to(device) for x in x_list]

        # Get latent representation (mu)
        mu, _ = model.encode(x_list)

        # Store latent vectors and labels
        latent_vectors.append(mu.cpu().numpy())
        labels.append(y.cpu().numpy())

# Concatenate batches
latent_vectors = np.concatenate(latent_vectors, axis=0)
labels = np.concatenate(labels, axis=0)

# Apply dimensionality reduction
if plot_type == 'tsne':
    reducer = TSNE(n_components=2, random_state=42)
    reduced_data = reducer.fit_transform(latent_vectors)
    title = 't-SNE Visualization of Latent Space'
else: # pca
    from sklearn.decomposition import PCA
    reducer = PCA(n_components=2, random_state=42)
    reduced_data = reducer.fit_transform(latent_vectors)
    title = 'PCA Visualization of Latent Space'

# Plot
plt.figure(figsize=(10, 8))
scatter = plt.scatter(reduced_data[:, 0], reduced_data[:, 1], c=labels, alpha=0.6,
                      cmap='coolwarm', edgecolors='w', s=100)

plt.colorbar(scatter, label='Complication')
plt.title(title)
plt.xlabel('Dimension 1')
plt.ylabel('Dimension 2')
plt.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.show()

def visualize_group_sparsity(model):
    """
    Visualize the group sparsity patterns

    Args:
        model: Trained model
    """
    # Get group sparsity matrix
    norm_matrix = model.get_group_sparsity_visualization()

    plt.figure(figsize=(12, 8))
    sns.heatmap(norm_matrix, cmap='viridis',
                xticklabels=[f'z{i+1}' for i in range(model.latent_dim)],
                yticklabels=model.timepoint_names,
                annot=False, cbar=True)

    plt.title('Group Sparsity Pattern')
    plt.xlabel('Latent Dimensions')
    plt.ylabel('Timepoints')
    plt.tight_layout()
    plt.show()

    # Analyze which latent dimensions are active for each timepoint
    for i, tp in enumerate(model.timepoint_names):
        active_dims = np.where(norm_matrix[i] > 0.1)[0] # Threshold can be adjusted
        print(f"Timepoint {tp} is influenced by latent dimensions: {list(active_dims + 1)}")

def visualize_attribute_regularization(model, test_loader, device):
    """
    Visualize how latent dimensions correspond to clinical attributes

    Args:
        model: Trained model
        test_loader: DataLoader for test data
        device: Device to use
    """
    model.eval()
    latent_vectors = []
    all_attributes = []

```

```

with torch.no_grad():
    for x_list, _, attributes in test_loader:
        # Move data to device
        x_list = [x.to(device) for x in x_list]

        # Get latent representation (mu)
        mu, _ = model.encode(x_list)

        # Store latent vectors and attributes
        latent_vectors.append(mu.cpu().numpy())
        all_attributes.append(attributes.cpu().numpy())

# Concatenate batches
latent_vectors = np.concatenate(latent_vectors, axis=0)
all_attributes = np.concatenate(all_attributes, axis=0)

# Calculate correlations between latent dimensions and attributes
corr_matrix = np.zeros((model.latent_dim, all_attributes.shape[1]))

for i in range(model.latent_dim):
    for j in range(all_attributes.shape[1]):
        corr = np.corrcoef(latent_vectors[:, i], all_attributes[:, j])[0, 1]
        corr_matrix[i, j] = corr

# Plot correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, cmap='coolwarm', center=0,
            xticklabels=list(model.attribute_dims.keys()),
            yticklabels=[f'z{i+1}' for i in range(model.latent_dim)],
            annot=True, fmt='.2f')

plt.title('Correlation between Latent Dimensions and Clinical Attributes')
plt.xlabel('Clinical Attributes')
plt.ylabel('Latent Dimensions')
plt.tight_layout()
plt.show()

# Print most correlated attributes for each latent dimension
for i in range(model.latent_dim):
    top_attr_idx = np.abs(corr_matrix[i]).argsort()[-3:][:-1] # Top 3
    top_attrs = [(list(model.attribute_dims.keys())[j], corr_matrix[i, j]) for j in top_attr_idx]
    print(f"Latent dimension z{i+1} is most correlated with: {top_attrs}")

```

```

def evaluate_model(model, test_loader, device):
    """
    Evaluate the model on test data

    Args:
        model: Trained model
        test_loader: DataLoader for test data
        device: Device to use

    Returns:
        metrics: Dictionary of test metrics
    """
    # Get test metrics
    test_metrics = model.trainer.validate(test_loader)

    # Print test metrics
    print("Test metrics:")
    for k, v in test_metrics.items():
        print(f"  {k}: {v:.4f}")

    # Calculate predictions for confusion matrix
    model.eval()
    all_preds = []
    all_targets = []

    with torch.no_grad():
        for x_list, y, attributes in test_loader:
            # Move data to device
            x_list = [x.to(device) for x in x_list]

            # Forward pass
            outputs = model(x_list, attributes)
            y_pred = outputs[4]

            # Apply sigmoid and threshold
            y_pred = torch.sigmoid(y_pred).cpu().numpy() > 0.5

            # Store predictions and targets
            all_preds.append(y_pred)
            all_targets.append(y.cpu().numpy())

    # Concatenate batches

```

```

all_preds = np.concatenate(all_preds, axis=0).flatten()
all_targets = np.concatenate(all_targets, axis=0).flatten()

# Plot confusion matrix
cm = confusion_matrix(all_targets, all_preds)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Complication', 'Complication'],
            yticklabels=['No Complication', 'Complication'])

plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.tight_layout()
plt.show()

return test_metrics

```

```

# Set up TensorBoard
from torch.utils.tensorboard import SummaryWriter
import datetime

def run_experiment(config=None):
    """
    Run a complete training experiment

    Args:
        config: Dictionary of configuration parameters

    Returns:
        model: Trained model
        history: Training history
    """
    if config is None:
        config = {
            # Model parameters
            'latent_dim': 32,
            'embed_dim': 16,
            'hidden_dim': 64,
            'medical_safety': True,

            # Training parameters

```

```

        'batch_size': 32,
        'lr': 1e-3,
        'weight_decay': 1e-5,
        'n_epochs': 100,
        'early_stopping_patience': 10,

        # Loss weights
        'beta': 1.0,      # KL weight
        'gamma': 0.1,    # Attribute regularization weight
        'delta': 0.1,    # Medical safety weight
        'lambda_reg': 0.5, # Group sparsity weight
    }

# Create output directory
timestamp = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
output_dir = f"./runs/{timestamp}"
os.makedirs(output_dir, exist_ok=True)

# Save configuration
with open(os.path.join(output_dir, 'config.json'), 'w') as f:
    json.dump(config, f, indent=4)

# Create TensorBoard writer
writer = SummaryWriter(output_dir)

# Create TensorBoard callback
tb_callback = TensorboardCallback(writer)

# Load data
print("Loading data...")
X_timepoints, y, attributes = load_mi_data()
train_loader, val_loader, test_loader = create_data_loaders(
    X_timepoints, y, attributes,
    batch_size=config['batch_size']
)

# Get input dimensions and timepoint names
input_dims = [X.shape[1] for X in X_timepoints.values()]
timepoint_names = list(X_timepoints.keys())
attribute_dims = {attr: i for i, attr in enumerate(attributes.columns)}

print(f"Input dimensions: {input_dims}")

```

```

print(f"Timepoint names: {timepoint_names}")
print(f"Attribute dimensions: {attribute_dims}")

# Initialize model
print("Initializing model...")
model = GroupInterpretableVAE(
    input_dims=input_dims,
    timepoint_names=timepoint_names,
    latent_dim=config['latent_dim'],
    embed_dim=config['embed_dim'],
    hidden_dim=config['hidden_dim'],
    attribute_dims=attribute_dims,
    medical_safety=config['medical_safety']
)

# Initialize trainer
trainer = GroupInterpretableVAETrainer(
    model=model,
    train_loader=train_loader,
    val_loader=val_loader,
    device=device,
    lr=config['lr'],
    weight_decay=config['weight_decay'],
    beta=config['beta'],
    gamma=config['gamma'],
    delta=config['delta'],
    lambda_reg=config['lambda_reg'],
    callbacks=[tb_callback]
)

# Train model
print("Training model...")
history, model = trainer.train(
    n_epochs=config['n_epochs'],
    early_stopping_patience=config['early_stopping_patience'],
    checkpoint_dir=os.path.join(output_dir, 'checkpoints')
)

# Attach trainer to model for later evaluation
model.trainer = trainer

# Save final model

```

```

torch.save({
    'model_state_dict': model.state_dict(),
    'model_config': {
        'input_dims': input_dims,
        'timepoint_names': timepoint_names,
        'latent_dim': config['latent_dim'],
        'embed_dim': config['embed_dim'],
        'hidden_dim': config['hidden_dim'],
        'attribute_dims': attribute_dims,
        'medical_safety': config['medical_safety']
    }
}, os.path.join(output_dir, 'final_model.pt'))

# Plot training history
plot_training_history(history)

# Visualize latent space
visualize_latent_space(model, test_loader, device)

# Visualize group sparsity
visualize_group_sparsity(model)

# Visualize attribute regularization
visualize_attribute_regularization(model, test_loader, device)

# Evaluate model on test data
test_metrics = evaluate_model(model, test_loader, device)

# Save test metrics
with open(os.path.join(output_dir, 'test_metrics.json'), 'w') as f:
    # Convert values to float for JSON serialization
    test_metrics_json = {k: float(v) for k, v in test_metrics.items()}
    json.dump(test_metrics_json, f, indent=4)

print(f"Experiment complete. Results saved to {output_dir}")
print(f"View TensorBoard with: tensorboard --logdir={output_dir}")

return model, history

# Run the experiment
model, history = run_experiment()

```



```

Loading data...
Loading MI Complications dataset from UCI repository...
Dataset shape: (1700, 111)
Number of samples: 1700
Number of complications: 12
Positive samples (with complications): 1037 (61.00%)
Timepoint admission: 3 features
Timepoint 24h: 0 features
Timepoint 48h: 0 features
Timepoint 72h: 0 features
Selected 6 clinical attributes: ['AGE', 'S_AD_ORIT', 'D_AD_ORIT', 'K_BLOOD', 'L_BLOOD', 'TIM
Timepoint admission: 2.63% missing values
Data preprocessing completed successfully.
Creating train/val/test splits...
Train set: 1190 samples (61.01% positive)
Validation set: 340 samples (60.88% positive)
Test set: 170 samples (61.18% positive)
Data loaders created successfully.
Input dimensions: [3]
Timepoint names: ['admission']
Attribute dimensions: {'AGE': 0, 'S_AD_ORIT': 1, 'D_AD_ORIT': 2, 'K_BLOOD': 3, 'L_BLOOD': 4,
Initializing model...
Training model...
Starting training for 100 epochs with early stopping patience 10

Epoch 1 [Train]:   0%|               | 0/38 [00:00<?, ?it/s]

Validation:   0%|               | 0/11 [00:00<?, ?it/s]

Epoch 1/100 - Train Loss: 1.6367, Val Loss: 1.6965, Train AUROC: 0.5211, Val AUROC: 0.6186
Saved best model at epoch 1 with val_loss: 1.6965

Epoch 2 [Train]:   0%|               | 0/38 [00:00<?, ?it/s]

Validation:   0%|               | 0/11 [00:00<?, ?it/s]

Epoch 2/100 - Train Loss: 1.2683, Val Loss: 1.2480, Train AUROC: 0.5875, Val AUROC: 0.6445
Saved best model at epoch 2 with val_loss: 1.2480

Epoch 3 [Train]:   0%|               | 0/38 [00:00<?, ?it/s]

```

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 3/100 - Train Loss: 1.1792, Val Loss: 1.1448, Train AUROC: 0.5971, Val AUROC: 0.6428
Saved best model at epoch 3 with val_loss: 1.1448

Epoch 4 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 4/100 - Train Loss: 1.1318, Val Loss: 1.0896, Train AUROC: 0.6094, Val AUROC: 0.6348
Saved best model at epoch 4 with val_loss: 1.0896

Epoch 5 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 5/100 - Train Loss: 1.1144, Val Loss: 1.0714, Train AUROC: 0.6288, Val AUROC: 0.6145
Saved best model at epoch 5 with val_loss: 1.0714

Epoch 6 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 6/100 - Train Loss: 1.1255, Val Loss: 1.0729, Train AUROC: 0.6115, Val AUROC: 0.6226

Epoch 7 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 7/100 - Train Loss: 1.1185, Val Loss: 1.0636, Train AUROC: 0.6174, Val AUROC: 0.6089
Saved best model at epoch 7 with val_loss: 1.0636

Epoch 8 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 8/100 - Train Loss: 1.0961, Val Loss: 1.0410, Train AUROC: 0.6380, Val AUROC: 0.6618
Saved best model at epoch 8 with val_loss: 1.0410

Epoch 9 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 9/100 - Train Loss: 1.0927, Val Loss: 1.0749, Train AUROC: 0.6234, Val AUROC: 0.6496

Epoch 10 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 10/100 - Train Loss: 1.0792, Val Loss: 1.0495, Train AUROC: 0.6407, Val AUROC: 0.6318

Epoch 11 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 11/100 - Train Loss: 1.0790, Val Loss: 1.0342, Train AUROC: 0.6354, Val AUROC: 0.6245
Saved best model at epoch 11 with val_loss: 1.0342

Epoch 12 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 12/100 - Train Loss: 1.0934, Val Loss: 1.0312, Train AUROC: 0.6357, Val AUROC: 0.6635
Saved best model at epoch 12 with val_loss: 1.0312

Epoch 13 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 13/100 - Train Loss: 1.0678, Val Loss: 1.0098, Train AUROC: 0.6256, Val AUROC: 0.6536
Saved best model at epoch 13 with val_loss: 1.0098

Epoch 14 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 14/100 - Train Loss: 1.0550, Val Loss: 1.0305, Train AUROC: 0.6453, Val AUROC: 0.6290

Epoch 15 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 15/100 - Train Loss: 1.0565, Val Loss: 1.0327, Train AUROC: 0.6461, Val AUROC: 0.6232

Epoch 16 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 16/100 - Train Loss: 1.0460, Val Loss: 1.0177, Train AUROC: 0.6584, Val AUROC: 0.6479

Epoch 17 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 17/100 - Train Loss: 1.0461, Val Loss: 1.0187, Train AUROC: 0.6455, Val AUROC: 0.6381

Epoch 18 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 18/100 - Train Loss: 1.0714, Val Loss: 1.0298, Train AUROC: 0.6224, Val AUROC: 0.6425

Epoch 19 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 19/100 - Train Loss: 1.0434, Val Loss: 1.0149, Train AUROC: 0.6553, Val AUROC: 0.6344

Epoch 20 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 20/100 - Train Loss: 1.0358, Val Loss: 1.0342, Train AUROC: 0.6456, Val AUROC: 0.6225

Epoch 21 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 21/100 - Train Loss: 1.0477, Val Loss: 0.9856, Train AUROC: 0.6273, Val AUROC: 0.6268
Saved best model at epoch 21 with val_loss: 0.9856

Epoch 22 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 22/100 - Train Loss: 1.0513, Val Loss: 1.0012, Train AUROC: 0.6206, Val AUROC: 0.6301

Epoch 23 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 23/100 - Train Loss: 1.0536, Val Loss: 1.0176, Train AUROC: 0.6470, Val AUROC: 0.6500

Epoch 24 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 24/100 - Train Loss: 1.0395, Val Loss: 1.0015, Train AUROC: 0.6568, Val AUROC: 0.6477

Epoch 25 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 25/100 - Train Loss: 1.0517, Val Loss: 1.0303, Train AUROC: 0.6445, Val AUROC: 0.6523

Epoch 26 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 26/100 - Train Loss: 1.0310, Val Loss: 1.0154, Train AUROC: 0.6313, Val AUROC: 0.6335

Epoch 27 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 27/100 - Train Loss: 1.0215, Val Loss: 0.9876, Train AUROC: 0.6566, Val AUROC: 0.6350

Epoch 28 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 28/100 - Train Loss: 1.0168, Val Loss: 1.0195, Train AUROC: 0.6352, Val AUROC: 0.6396

Epoch 29 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 29/100 - Train Loss: 1.0269, Val Loss: 1.0197, Train AUROC: 0.6426, Val AUROC: 0.6370

Epoch 30 [Train]: 0%| | 0/38 [00:00<?, ?it/s]

Validation: 0%| | 0/11 [00:00<?, ?it/s]

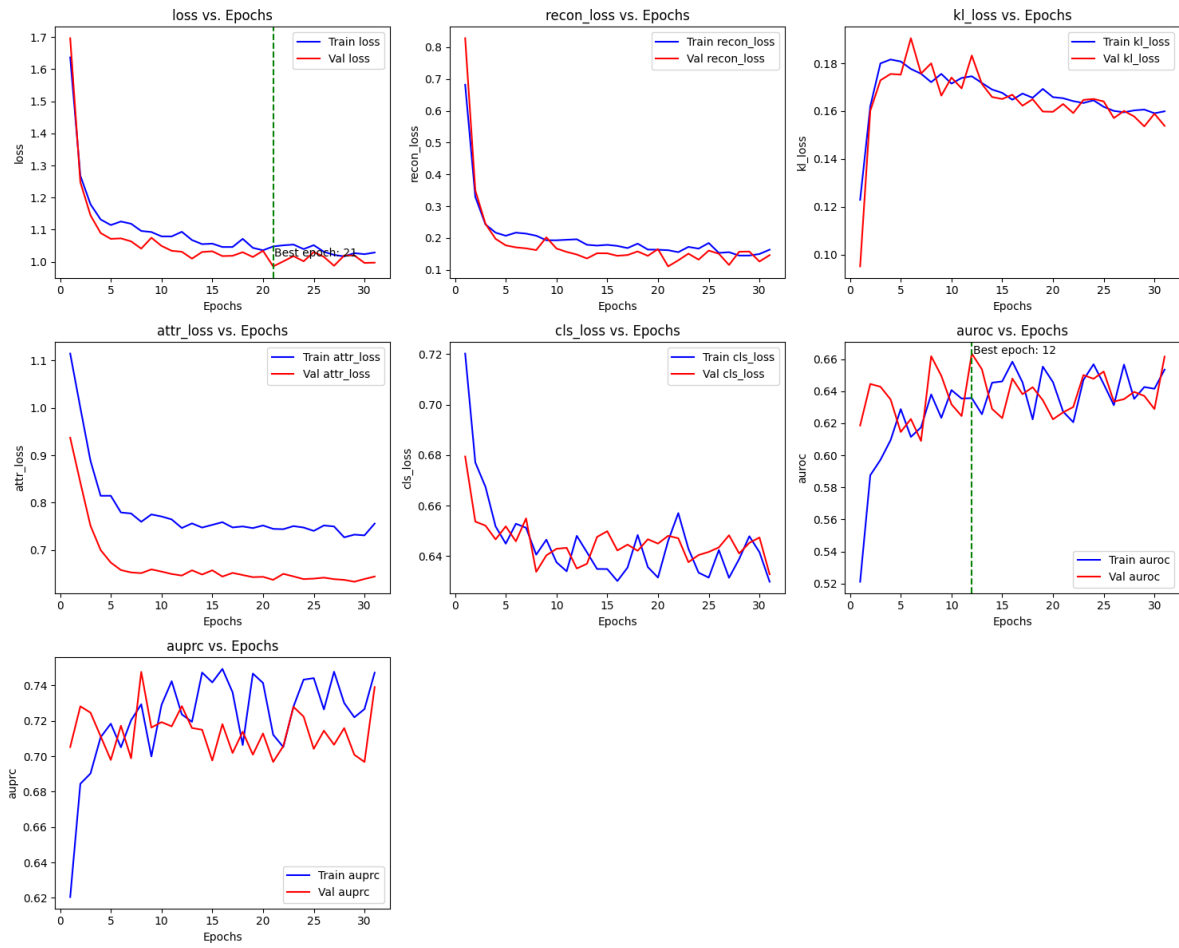
Epoch 30/100 - Train Loss: 1.0237, Val Loss: 0.9964, Train AUROC: 0.6415, Val AUROC: 0.6289

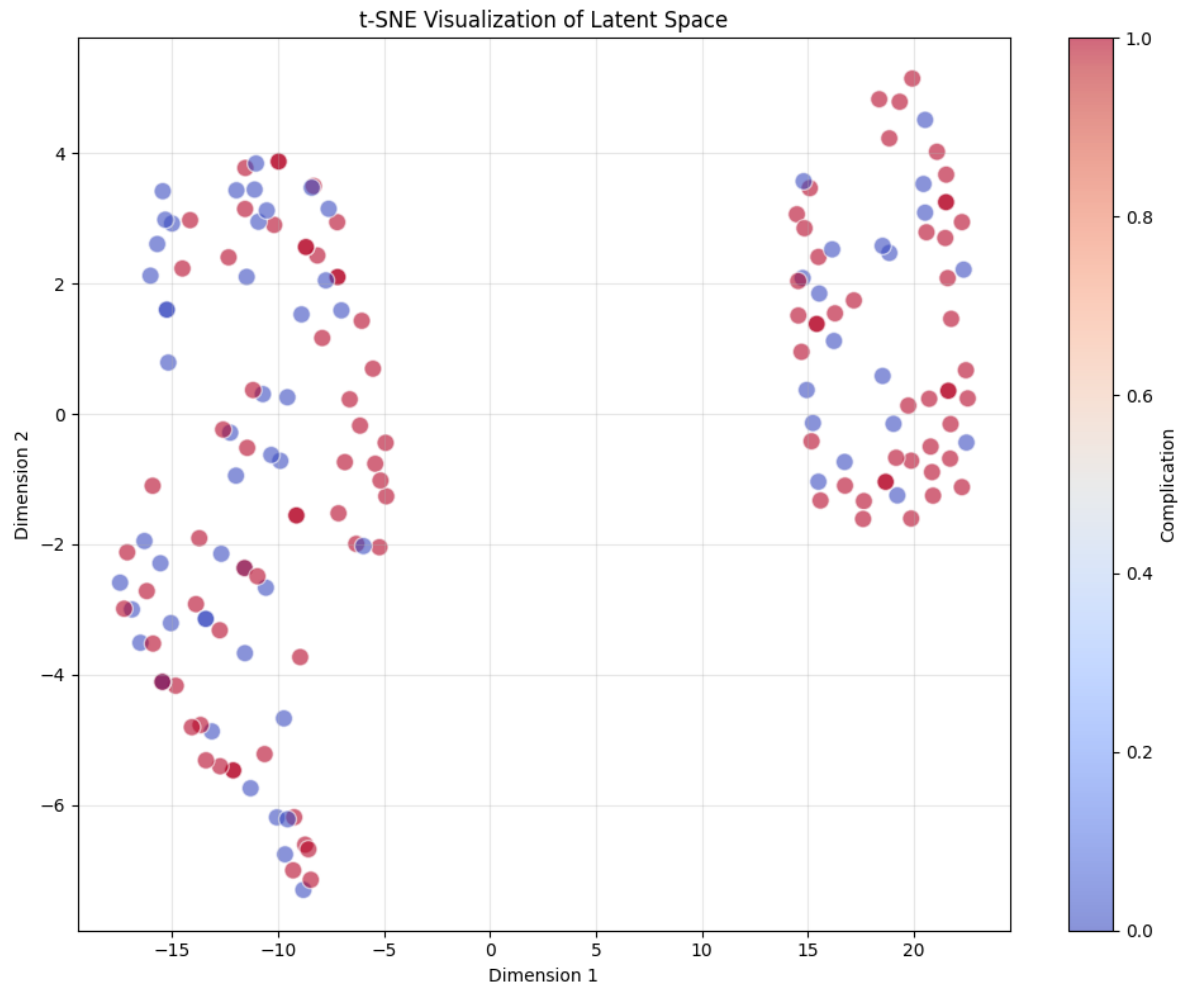
Epoch 31 [Train]: 0%| | 0/38 [00:00<?, ?it/s]


Validation: 0%| | 0/11 [00:00<?, ?it/s]

Epoch 31/100 - Train Loss: 1.0288, Val Loss: 0.9974, Train AUROC: 0.6535, Val AUROC: 0.6616


Early stopping triggered after 31 epochs







Timepoint admission is influenced by latent dimensions: []



Latent dimension z1 is most correlated with: [('TIME_B_S', np.float64(-0.6377436265755303)),

Latent dimension z2 is most correlated with: [('AGE', np.float64(-0.41800312911144777)), ('T

Latent dimension z3 is most correlated with: [('S_AD_ORIT', np.float64(-0.1842425731977365))

Latent dimension z4 is most correlated with: [('AGE', np.float64(0.6342007109107097)), ('TIM

Latent dimension z5 is most correlated with: [('TIME_B_S', np.float64(-0.34579389919543563))

Latent dimension z6 is most correlated with: [('TIME_B_S', np.float64(-0.7280939836684573)),

Latent dimension z7 is most correlated with: [('AGE', np.float64(-0.313516619442179)), ('K_BI

Latent dimension z8 is most correlated with: [('TIME_B_S', np.float64(0.45559995667911973)),

Latent dimension z9 is most correlated with: [('AGE', np.float64(-0.3431252115532879)), ('TI

Latent dimension z10 is most correlated with: [('AGE', np.float64(0.8897627722811079)), ('TI

Latent dimension z11 is most correlated with: [('TIME_B_S', np.float64(-0.6517921717200551))

Latent dimension z12 is most correlated with: [('AGE', np.float64(0.9294212512777801)), ('TI

Latent dimension z13 is most correlated with: [('TIME_B_S', np.float64(0.9474185916893575)),

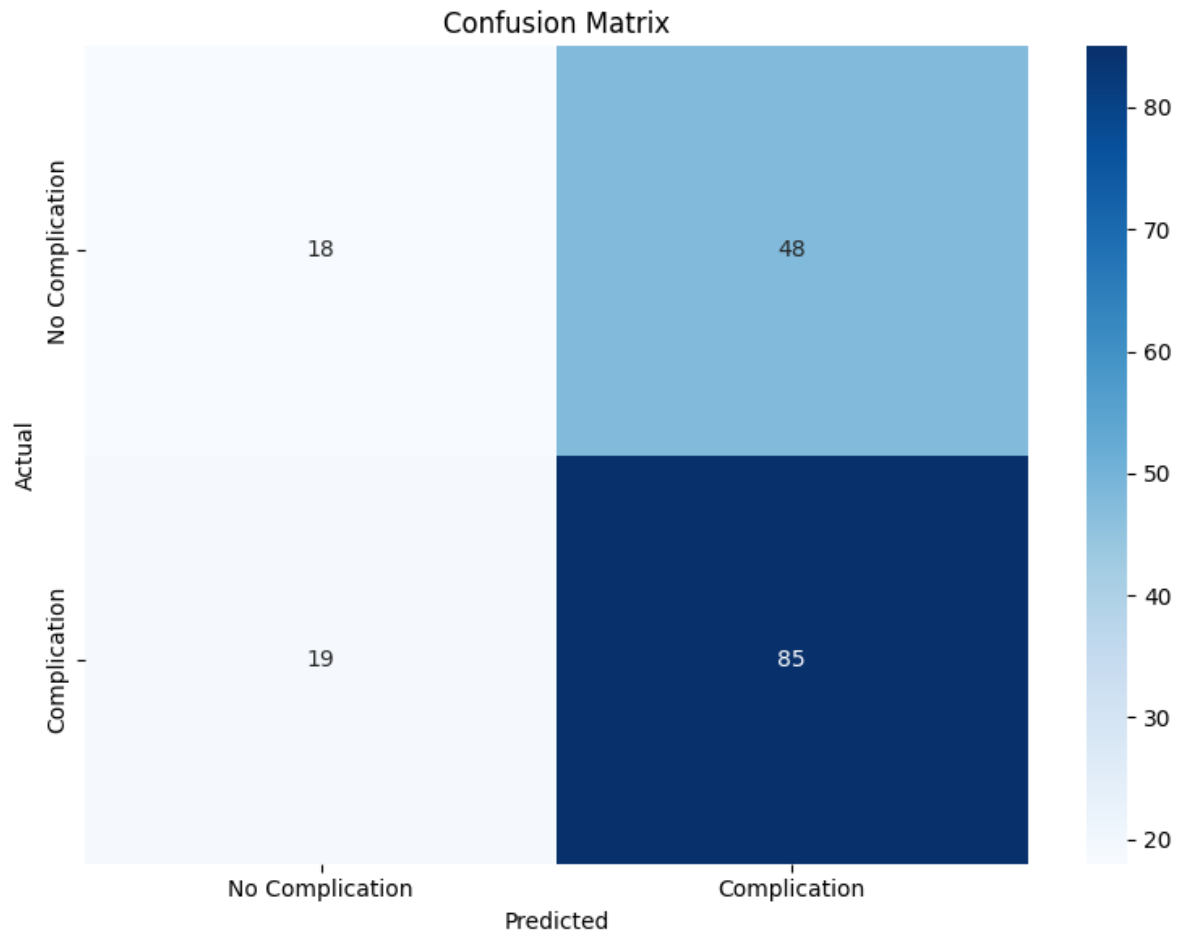
Latent dimension z14 is most correlated with: [('TIME_B_S', np.float64(-0.7869936919339381))

Latent dimension z15 is most correlated with: [('TIME_B_S', np.float64(-0.6971735505807037))]
Latent dimension z16 is most correlated with: [('S_AD_ORIT', np.float64(-0.16445866893025324))]
Latent dimension z17 is most correlated with: [('AGE', np.float64(0.3243954744348347)), ('TIME_B_S', np.float64(0.7449000791157308))]
Latent dimension z18 is most correlated with: [('TIME_B_S', np.float64(0.7449000791157308))]
Latent dimension z19 is most correlated with: [('AGE', np.float64(-0.8620416975866418)), ('TIME_B_S', np.float64(0.6177065356762959))]
Latent dimension z20 is most correlated with: [('AGE', np.float64(0.6177065356762959)), ('TIME_B_S', np.float64(0.5387503236709416))]
Latent dimension z21 is most correlated with: [('AGE', np.float64(0.5387503236709416)), ('S_AD_ORIT', np.float64(-0.24146180427697553))]
Latent dimension z22 is most correlated with: [('D_AD_ORIT', np.float64(-0.24146180427697553))]
Latent dimension z23 is most correlated with: [('TIME_B_S', np.float64(0.5080572993263404))]
Latent dimension z24 is most correlated with: [('AGE', np.float64(0.26538603210319356)), ('S_AD_ORIT', np.float64(0.4620892951812275))]
Latent dimension z25 is most correlated with: [('TIME_B_S', np.float64(0.4620892951812275))]
Latent dimension z26 is most correlated with: [('AGE', np.float64(0.8740096872569565)), ('TIME_B_S', np.float64(-0.4716705271016905))]
Latent dimension z27 is most correlated with: [('TIME_B_S', np.float64(-0.4716705271016905))]
Latent dimension z28 is most correlated with: [('TIME_B_S', np.float64(0.36638965051109923))]
Latent dimension z29 is most correlated with: [('TIME_B_S', np.float64(-0.47575436629615836))]
Latent dimension z30 is most correlated with: [('TIME_B_S', np.float64(-0.4748258954726793))]
Latent dimension z31 is most correlated with: [('TIME_B_S', np.float64(0.6256528170199492))]
Latent dimension z32 is most correlated with: [('TIME_B_S', np.float64(-0.7072562966673279))]

Validation: 0% | 0/6 [00:00<?, ?it/s]

Test metrics:

loss: 1.0215
recon_loss: 0.1466
kl_loss: 0.1715
attr_loss: 0.7239
cls_loss: 0.6310
auROC: 0.6719
auprc: 0.7851



Experiment complete. Results saved to ./runs/20250424-025550
View TensorBoard with: tensorboard --logdir=./runs/20250424-025550